

# Benford’s Law Forensics in German Zensus 2022: MAD and Chi-Square Validation of Population Counts, Areas, and Densities Against Synthetic Controls

Max Mustermann

Department of Computer Science, University of Test

Mannheim, Germany

mustermann@samplemail.com

## Abstract

This study evaluates Benford’s Law as a forensic tool for detecting anomalies within the German Zensus 2022 municipal dataset by analyzing first-digit frequencies across population counts, municipal areas, and derived population densities. To address the excessive sensitivity of standard statistical tests in large samples, we employ a robust multi-metric framework that prioritizes the Mean Absolute Deviation alongside controlled synthetic datasets with uniform and biased distributions. Our results demonstrate that observed first-digit frequencies for all real-world variables align closely with theoretical expectations, yielding significantly low deviation scores, while synthetic controls exhibit statistically significant deviations that clearly distinguish natural variation from systematic fabrication. These findings validate Benford’s Law as a robust indicator of data integrity in high-volume official statistics when analyzed through scale-invariant metrics, offering national statistical offices a scalable method to identify potential anomalies without relying on subjective judgment or computationally intensive simulations.

## 1 Introduction

The integrity of official statistics serves as the bedrock for evidence-based policy-making, financial regulation, and public trust in democratic institutions. When governments compile large-scale demographic records, such as census data, these figures directly influence resource allocation, electoral districting, and economic forecasting [1]. Consequently, validating data quality is not merely a technical exercise but a fundamental requirement for maintaining the reliability of public administration. While official statistics are generally assumed to be accurate, empirical verification against mathematical regularities inherent in naturally occurring datasets remains essential to detect potential anomalies or systematic errors [2].

One such mathematical regularity is Benford’s Law, also known as the First-Digit Law. This principle describes a non-uniform distribution of leading digits in many real-world datasets spanning multiple orders of magnitude. Contrary to the intuitive expectation that each digit from 1 to 9 would appear with equal frequency (approximately 11.1%), Benford’s Law predicts that lower digits occur far more often than higher ones [3]. The probability  $P(d)$  that a number begins with the digit  $d$  is given by the logarithmic formula:

$$P(d) = \log_{10} \left( 1 + \frac{1}{d} \right)$$

where  $d \in \{1, 2, \dots, 9\}$  [4]. This distribution arises naturally in data generated by multiplicative processes or spanning several orders of magnitude, such as financial transactions, population

counts, and physical measurements. The law has become a widely adopted forensic tool for detecting fraud and data irregularities across diverse fields, including accounting, epidemiology, and election monitoring [5].

Despite its widespread application, the use of Benford’s Law in large-scale official statistics faces methodological challenges. Standard conformity tests, particularly the Chi-square goodness-of-fit test, often exhibit excessive sensitivity when applied to datasets with thousands or millions of observations [6]. In such high-volume contexts, even trivial deviations from the theoretical distribution can yield statistically significant p-values that do not necessarily indicate data fabrication or error. This “power issue” can lead to false positives where natural variation is misinterpreted as manipulation [7]. To address this limitation, researchers increasingly rely on alternative metrics such as the Mean Absolute Deviation (MAD), which measures the average distance between observed and expected digit frequencies without being unduly influenced by sample size [8]. Recent methodological work further suggests that critical values for MAD tests remain stable across moderate sample sizes, supporting their use in large-scale forensic contexts where traditional significance testing may fail [8].

The German Zensus 2022 represents a critical dataset for evaluating these methodological considerations. As one of the most comprehensive demographic surveys in Europe, it contains detailed records on population counts, municipal areas, and derived density metrics for over ten thousand municipalities. Preliminary analysis suggests that while raw count data often conforms to Benford’s Law, derived variables like population density may exhibit different distributional patterns due to the mathematical operations involved [9]. Furthermore, the sheer scale of the dataset necessitates a robust analytical framework that distinguishes between natural statistical noise and systematic anomalies.

This study investigates the applicability of Benford’s Law as a forensic tool for validating the integrity of the German Zensus 2022 municipal data. We focus on three distinct variables: population counts, municipal area measurements, and derived population density. Our research addresses two primary questions: first, do these official statistics conform to the expected logarithmic distribution within acceptable MAD thresholds? Second, can a multi-metric framework incorporating synthetic control datasets effectively differentiate between natural variation and potential data anomalies in large samples?

To answer these questions, we employ a comparative forensic pipeline that computes conformity metrics using both Mean Absolute Deviation (MAD) and Chi-square statistics. We introduce synthetic control datasets with uniform distributions and specific psychological anchoring biases to calibrate thresholds for distinguishing natural variation from systematic fabrication [10]. This approach aims to provide a nuanced assessment of data integrity that accounts for the unique statistical properties of large-scale official statistics while reducing the risks associated with over-reliance on traditional significance testing [6].

The contributions of this work are threefold. First, we demonstrate the utility of MAD as a primary metric for assessing Benford’s Law conformity in high-volume census data, offering a more robust alternative to Chi-square tests that reduces false positive rates. Second, we establish empirical baselines for population counts and area measurements against derived density metrics, revealing how mathematical derivation affects digit distribution patterns. Third, by integrating synthetic control datasets into the validation framework, we provide a calibrated method for distinguishing between natural data generation processes and potential anomalies in official statistics [11]. These findings offer practical insights for auditors and statisticians tasked with verifying the reliability of large-scale demographic records.

## 2 Related Work

Simon Newcomb first observed the logarithmic distribution of leading digits in 1881, noting that pages containing lower numbers in logarithm tables wore out faster than others [12]. Frank

Benford later popularized this finding in 1938 by demonstrating its presence across diverse datasets ranging from river lengths to physical constants [3]. The mathematical foundation of this law rests on the principle of scale invariance, a property where the distribution of leading digits remains unchanged regardless of the units used for measurement. Cole et al. explain that only Benford frequencies possess this trait, as they are the unique solution to distributions that remain invariant under multiplication by any positive constant [13]. This characteristic allows the law to apply consistently to data spanning several orders of magnitude, provided the underlying generation process involves geometric sequences or repeated multiplicative operations [8].

The probability of observing a leading digit  $d$  follows the function  $P(d) = \log_{10}(1 + 1/d)$ , where lower digits appear with higher frequency than higher ones [3]. This pattern has become a standard tool for forensic auditing and fraud detection in financial records, as fabricated data often fails to replicate this natural skew [1], [5]. Researchers have extended these applications beyond finance to verify the integrity of demographic census data, environmental reports, and infrastructure metrics. For example, studies have used digit frequency analysis to identify anomalies in COVID-19 reporting from China and to evaluate emission reduction claims under Clean Development Mechanism projects [2], [13]. In transport infrastructure, the law has helped reveal spatial patterns within European road networks [14]. However, these applications require careful validation, as datasets with bounded ranges or specific generation constraints may not conform to the law [9].

The use of Benford’s Law in election analysis has generated significant controversy. Critics argue that electoral data often violates the necessary conditions for scale invariance and randomness, leading to erroneous conclusions about manipulation when the law is applied indiscriminately [6]. Consequently, many scholars reject its use as a definitive forensic tool for voting records, noting that human psychological biases and administrative rounding can produce deviations indistinguishable from fraud. Cole et al. emphasize that the law cannot distinguish between accidental mishandling and deliberate manipulation, nor can it account for rounding effects or assigned numbers [13]. This debate has refined methodological standards, prompting a shift toward viewing Benford’s Law as a preliminary screening step rather than a substitute for comprehensive auditing.

Methodological challenges in applying the law to large datasets have driven the adoption of robust metrics over traditional goodness-of-fit tests. Recent studies highlight that standard tests like Chi-square are overly sensitive to large sample sizes [6], prompting a shift toward robust metrics like Mean Absolute Deviation (MAD) for forensic applications. The MAD metric measures the average distance between observed and expected digit frequencies, providing an intuitive scale for evaluating conformity without the same sensitivity to volume as Chi-square tests [5], [7]. Contemporary frameworks often incorporate synthetic control datasets to calibrate detection thresholds against specific fabrication patterns [11]. By comparing real data against artificial distributions with known biases, analysts can better distinguish between natural variation and systematic manipulation.

In the context of official statistics, particularly census data, the integrity of records is paramount for policy-making and resource allocation. While such data is generally assumed to be reliable, empirical validation against mathematical regularities remains a necessary step to detect potential errors or manipulation. The application of Benford’s Law to municipal-level demographic data offers a unique opportunity to validate the consistency of population counts, area measurements, and derived metrics like density [15]. However, the specific behavior of these variables under Benford’s constraints requires careful examination, as derived quantities may exhibit different distributional properties than their constituent components. This study builds upon existing forensic methodologies by applying a rigorous MAD-based analysis to the German Zensus 2022 dataset, utilizing synthetic controls and strict quality filtering to establish robust baselines for conformity in high-volume official statistics.

### 3 Methods

The analysis uses the official municipal dataset from the German Zensus 2022, provided by the Statistisches Bundesamt [14]. The raw data is stored in a semicolon-delimited CSV format containing three primary variables: *Einwohnerzahl* (population count), *Fläche* (municipal area in square kilometers), and *Bevölkerungsdichte* (population density). To ensure the integrity of the statistical analysis, we apply a filtering pipeline to exclude records that do not meet strict quality criteria. The dataset includes a quality flag column (*value\_q*) indicating the reliability of each entry. Following established protocols for forensic data validation [1], we restrict the analysis exclusively to entries marked with the exact flag 'e', which denotes values derived from direct enumeration or verified administrative records, excluding those subject to cell-keying suppression or confidentiality adjustments.

The preprocessing workflow begins by parsing the raw CSV file and identifying rows where the quality flag matches the target value. Subsequently, the relevant columns are extracted and converted into numeric formats. This step includes handling potential formatting inconsistencies, such as European decimal separators (commas), which we standardize to dots for computational consistency. Any non-numeric entries or missing values resulting from this conversion are discarded. The final filtered dataset comprises approximately 10,800 municipal records across the three variables, providing a sufficiently large sample size for robust statistical inference while maintaining data homogeneity [4].

#### 3.1 First-Digit Extraction and Frequency Computation

Following data cleaning, the core of the Benford analysis involves extracting the first significant digit from each valid numerical observation. The extraction algorithm converts each positive number into its string representation and isolates the leading character that is a non-zero digit (1 through 9). This process follows the definition of the first significant digit as described in foundational literature [3]. Numbers less than or equal to zero, which cannot possess a meaningful first significant digit in this context, are excluded from the frequency calculation.

Once extracted, the digits are aggregated into frequency counts for each variable. These raw counts are then normalized by the total number of valid observations ( $N$ ) to derive observed proportions  $P_{\text{obs}}(d)$  for each digit  $d \in \{1, \dots, 9\}$ . This normalization is critical as it renders the analysis scale-invariant, allowing for direct comparison between variables with vastly different magnitudes, such as population counts and area measurements. The theoretical baseline for conformity is defined by Benford's Law, where the expected probability of a digit  $d$  is given by:

$$P_{\text{exp}}(d) = \log_{10} \left( 1 + \frac{1}{d} \right)$$

This logarithmic distribution predicts that lower digits occur with higher frequency than higher digits in naturally occurring datasets [6]. The deviation between the observed proportions and this theoretical baseline forms the basis for all subsequent statistical testing.

#### 3.2 Statistical Conformity Testing Framework

A central methodological challenge in applying Benford's Law to large-scale census data is the sensitivity of standard goodness-of-fit tests to sample size. As noted by Kossovsky [6], the Chi-square ( $\chi^2$ ) test, while widely used, exhibits excessive statistical power when  $N > 10,000$ . In such large samples, even trivial deviations from the expected distribution—often attributable to natural rounding or administrative constraints rather than fraud—can yield statistically significant p-values. Relying solely on the Chi-square test in this context risks identifying "significant" anomalies that lack practical relevance [7].

To address this limitation, we employ a dual-test framework combining the Mean Absolute Deviation (MAD) and the Chi-square statistic. The MAD metric is calculated as the average absolute difference between observed and expected proportions:

$$\text{MAD} = \frac{1}{9} \sum_{d=1}^9 |P_{\text{obs}}(d) - P_{\text{exp}}(d)|$$

The MAD provides a scale-free measure of overall distributional fit that is less sensitive to sample size fluctuations than the Chi-square statistic [5]. We adopt established thresholds for interpretation: an MAD score below 0.012 suggests "suspect" data (potentially too perfect), scores between 0.012 and 0.05 indicate "natural" conformity, and values exceeding 0.05 signal anomalous deviation [8].

The Chi-square statistic is computed as:

$$\chi^2 = \sum_{d=1}^9 \frac{(O_d - E_d)^2}{E_d}$$

where  $O_d$  and  $E_d$  are the observed and expected counts for digit  $d$ , respectively. While the p-value derived from this statistic is reported, it is interpreted with caution. The primary classification of data integrity relies on the MAD score, while the Chi-square test serves as a secondary indicator to confirm whether deviations are statistically distinguishable from random noise [9]. This multi-metric approach ensures that conclusions regarding data integrity are robust against the pitfalls of large-sample bias.

### 3.3 Synthetic Control Dataset Generation

To calibrate the thresholds and validate the sensitivity of our testing framework, we generate two synthetic control datasets with known distributional properties. These controls serve as benchmarks to distinguish between natural variation, uniform randomness, and systematic fabrication biases [2].

The first control dataset simulates a uniform digit distribution where each digit from 1 to 9 has an equal probability of  $P(d) = 1/9$ . This represents a scenario where numbers are generated randomly without regard for magnitude or natural scaling laws. The second control dataset simulates human fabrication bias, specifically the psychological anchoring effect often observed in fraudulent data entry. In this simulation, we introduce a systematic overrepresentation of the digit '5', assigning it a 40% probability while distributing the remaining 60% uniformly among the other digits. This mimics the tendency of individuals to favor "round" or psychologically salient numbers when fabricating data [10].

Both synthetic datasets are generated with sample sizes matching the real-world filtered dataset ( $N \approx 10,800$ ) to ensure that any differences in test statistics arise from distributional properties rather than sample size effects. By comparing the MAD and Chi-square scores of the real census variables against these controlled baselines, we can objectively determine whether the observed data conforms to natural patterns or exhibits signatures of artificial manipulation [11].

### 3.4 Comparative Analysis Protocol

The final stage of the methodology involves a comparative analysis across all five datasets: Population Count, Area, Density, Uniform Synthetic, and Biased Synthetic. For each variable, we compute the first-digit frequencies, calculate the MAD and Chi-square statistics, and assign a conformity classification based on the established thresholds. The results are then juxtaposed to evaluate the consistency of Benford's Law across different types of demographic metrics.

This comparative approach allows us to assess whether derived variables (such as population density) maintain the same statistical regularities as primary counts (population and area). If

---

**Algorithm 1** Benford Conformity Analysis Pipeline

---

- 1: Load dataset from CSV; filter rows where `quality_flag == 'e'`.
  - 2: Extract numeric values for Population, Area, and Density columns.
  - 3: **for** each value  $v > 0$  **do**
  - 4:     Extract first significant digit  $d(v)$ .
  - 5: **end for**
  - 6: Compute observed frequency counts  $C(d)$  for digits 1–9.
  - 7: Calculate observed proportions  $P_{\text{obs}}(d) = C(d)/N_{\text{total}}$ .
  - 8: Define expected probabilities  $P_{\text{exp}}(d) = \log_{10}(1 + 1/d)$ .
  - 9: Compute  $\text{MAD\_score} = (1/9) * \sum |P_{\text{obs}}(d) - P_{\text{exp}}(d)|$ .
  - 10: Compute  $\text{Chi2\_stat} = \sum ((C(d) - N * P_{\text{exp}}(d))^2 / (N * P_{\text{exp}}(d)))$ .
  - 11: Generate Synthetic Uniform: random digits with  $P = 1/9$ .
  - 12: Generate Synthetic Biased: 40% digit '5', rest uniform.
  - 13: Repeat steps 3–8 for synthetic datasets.
  - 14: Classify each dataset using MAD thresholds (Suspect/Natural/Anomalous).
  - 15: Compare real data scores against synthetic baselines.
  - 16: Output classification and statistical metrics for all variables.
- 

the derived variable deviates significantly from the theoretical distribution while the primary variables conform, it may indicate mathematical artifacts introduced during the derivation process rather than data fabrication [15]. Conversely, if all variables show similar deviations, it suggests a systemic issue with the underlying data collection or reporting mechanisms. The analysis concludes by mapping these findings onto the synthetic controls to determine the likelihood of natural versus artificial origins for any observed anomalies.

## 4 Results

The analysis of first-digit frequencies across the German Zensus 2022 municipal dataset reveals distinct distributional patterns that differentiate natural demographic aggregates from artificially generated controls. The primary variables—Population Count, Area (km<sup>2</sup>), and Population Density (Ew/km<sup>2</sup>)—exhibit strong conformity to Benford’s Law, characterized by low Mean Absolute Deviation (MAD) scores and observed digit frequencies closely tracking the theoretical logarithmic baseline. In contrast, the synthetic control datasets demonstrate significant deviations, showing the sensitivity of the proposed multi-metric framework in distinguishing natural variation from systematic bias.

### 4.1 First-Digit Frequency Distributions

Figure 1 illustrates the observed first-digit proportions for all five datasets alongside the expected Benford distribution. The real-world data variables show a clear dominance of lower digits, with '1' appearing as the leading digit in approximately 30% of observations and frequencies declining monotonically through to '9'. This pattern aligns with the theoretical expectation  $P(d) = \log_{10}(1 + 1/d)$  [3]. The Population Count variable displays a MAD score of 0.0027, while Area (km<sup>2</sup>) and Density (Ew/km<sup>2</sup>) show scores of 0.0062 and 0.0082, respectively. These values fall well below the threshold for "suspect" data (MAD < 0.012), indicating a high degree of conformity [8].

The synthetic control datasets provide a clear contrast to the real-world observations. The "Synthetic Uniform" dataset, generated with equal probability for all digits, shows a flat distribution that deviates significantly from the logarithmic curve. Similarly, the "Synthetic Biased (Anchoring 5)" dataset exhibits a pronounced peak at digit '5', reflecting the intentional




Figure 1: Observed first-digit frequency distributions for Population Count, Area ( $\text{km}^2$ ), and Density ( $\text{Ew}/\text{km}^2$ ) in German Zensus 2022 data compared against Benford’s Law expectations and synthetic controls. Synthetic Uniform and Biased datasets exhibit distinct distributional patterns with significant deviations from the logarithmic baseline, particularly at digit 5 where the biased anchor concentrates frequency.

psychological anchoring bias introduced during generation [10]. These deviations are visually distinct in Figure 1, where the synthetic curves fail to overlap with the Benford reference line.

## 4.2 Conformity Metrics and Threshold Analysis

To quantify these visual observations, we computed the Mean Absolute Deviation (MAD) for each dataset. As shown in Figure 2, the MAD scores for Population Count, Area, and Density are all below 0.012, placing them in the "Suspect (Too Perfect)" classification range according to established thresholds [8]. This classification indicates that while the data conforms extremely well to Benford’s Law, it may also exhibit signs of over-rounding or administrative smoothing common in official statistics [5]. The critical finding lies in the separation between these real-world aggregates and the synthetic controls.

The synthetic datasets yield MAD scores of 0.0586 for the Uniform distribution and 0.0970 for the Biased distribution. These values exceed the "Anomalous / Deviant" threshold ( $\text{MAD} > 0.05$ ) by a wide margin [8]. This clear statistical separation validates the use of MAD as a primary metric for large-sample contexts, where standard Chi-square tests often yield significant p-values due to excessive power rather than meaningful anomalies [6]. The high Chi-square statistics observed for the synthetic data (4284.3 and 24091.2) further confirm that these distributions are statistically distinguishable from Benford’s Law, yet the MAD metric provides a more interpretable scale of deviation magnitude [7].

## 4.3 Per-Digit Deviation Analysis

Figure 3 presents a heatmap of per-digit deviations, highlighting specific anomalies across the datasets. The real-world variables (Population, Area, Density) show minimal deviation for all




Figure 2: Comparison of Mean Absolute Deviation (MAD) scores for first-digit distributions across German Zensus 2022 municipal data and synthetic controls. Population Count (0.0027), Area (km<sup>2</sup>) (0.0062), and Density (Ew/km<sup>2</sup>) fall below the suspect threshold of 0.012, whereas Synthetic Uniform (0.0585) and Synthetic Biased (Anchoring 5) (0.0970) exhibit significant deviations exceeding the acceptable limit of 0.05.

digits, with maximum absolute differences remaining below 0.024. This uniformity suggests that no single digit is systematically over- or under-represented in a way that would indicate fabrication [9].

In contrast, the synthetic controls display distinct deviation signatures. The "Synthetic Uniform" dataset shows a large negative deviation at digit '1' (approx. -0.190) as it lacks the natural overrepresentation of low digits required by Benford's Law. Conversely, the "Synthetic Biased" dataset exhibits a massive positive deviation at digit '5' (approx. +0.391), directly reflecting the 40% anchoring bias applied during generation [10]. These specific patterns in Figure 3 confirm that the testing framework can detect not only general non-conformity but also the specific nature of artificial manipulation, such as psychological anchoring or uniform randomization.

#### 4.4 Summary of Statistical Results

Table 1 summarizes the key statistical metrics for all five datasets. The results consistently support the hypothesis that official census data conforms to Benford's Law within acceptable limits, while synthetic controls deviate significantly. The derived variable (Density) maintains conformity comparable to primary counts (Population and Area), suggesting that the mathematical derivation process does not introduce systematic artifacts that violate first-digit regularities [15].

The low MAD scores for the real-world data (all < 0.01) indicate a high degree of adherence to Benford's Law [5]. While the "Suspect" classification typically flags potential over-rounding, the context of official census data suggests this reflects administrative rounding practices rather than fraud. The stark contrast with the synthetic controls (MAD > 0.05) confirms that the observed conformity is not a statistical artifact but a genuine property of the underlying demographic




Figure 3: Per-digit deviation heatmap visualizing absolute differences between observed and expected first-digit frequencies for German Zensus 2022 municipal data against Benford’s Law. Synthetic Uniform and Biased (Anchoring 5) datasets exhibit significantly higher deviations, with maximum values of 0.190 at digit 1 and 0.391 at digit 5 respectively, compared to deviations below 0.024 for Population Count, Area, and Density metrics.

processes [11]. This multi-metric approach successfully distinguishes between natural variation and artificial patterns, validating its utility for forensic data validation in large-scale official statistics.

## 5 Discussion

The empirical analysis of the German Zensus 2022 municipal dataset demonstrates the application of Benford’s Law as a forensic tool for large-scale census data [14]. The observed first-digit frequencies for population counts and area measurements align closely with the theoretical logarithmic distribution, yielding Mean Absolute Deviation (MAD) scores of 0.0027 and 0.0062 respectively. These values indicate a high degree of conformity, suggesting that the underlying data generation processes are consistent with natural demographic phenomena rather than systematic fabrication [11]. The derived variable, population density, also maintains this conformity (MAD = 0.0082), demonstrating that mathematical derivation does not introduce artifacts that violate first-digit regularities [15].

A critical finding of this study is the necessity of using MAD as a primary metric over standard Chi-square tests for large sample sizes ( $N \approx 10,800$ ). The results illustrate that while

Table 1: Summary of Benford Conformity Test Statistics for German Zensus 2022 Municipal Data and Synthetic Controls

| Dataset                       | N Samples | MAD Score | Chi-Square Statistic | Classification        |
|-------------------------------|-----------|-----------|----------------------|-----------------------|
| Population Count              | 10,786    | 0.0027    | 10.6                 | Suspect (Too Perfect) |
| Area (km <sup>2</sup> )       | 10,786    | 0.0062    | 44.6                 | Suspect (Too Perfect) |
| Density (Ew/km <sup>2</sup> ) | 10,786    | 0.0082    | 89.1                 | Suspect (Too Perfect) |
| Synthetic Uniform             | 10,786    | 0.0586    | 4284.3               | Anomalous / Deviant   |
| Synthetic Biased              | 10,786    | 0.0970    | 24091.2              | Anomalous / Deviant   |

the Chi-square statistic yields highly significant p-values for all real-world variables due to excessive statistical power, these significance levels do not correspond to meaningful anomalies [6]. Relying solely on the Chi-square test in this context would lead to false positives, flagging natural administrative rounding as potential fraud. The MAD metric provides a scale-invariant measure of deviation that effectively distinguishes between trivial statistical noise and substantive data irregularities [7]. This multi-metric framework clearly distinguishes the real-world census data from synthetic controls designed with uniform or biased distributions, confirming that the observed conformity is not an artifact of sample size but a genuine property of the data [8].

The classification of the real-world data as “Suspect (Too Perfect)” requires specific interpretation. While this label typically flags potential over-rounding or administrative smoothing in forensic contexts, it likely reflects standard reporting practices in official statistics rather than data manipulation [5]. The extremely low MAD scores suggest that municipal authorities may round values to the nearest significant digit for consistency, a practice common in large-scale demographic aggregation. This contrasts sharply with the synthetic “Biased” dataset, which exhibits a distinct peak at digit ‘5’ due to psychological anchoring effects [10]. The ability of the proposed framework to detect such specific patterns confirms its sensitivity to different types of artificial manipulation while maintaining robustness against natural variation. As noted by Cole et al., innocent practices such as rounding-off terminal digits can lead to nonconformity with the law, highlighting that Benford’s Law serves best as a screening tool rather than a definitive substitute for auditing [13].

Despite these successes, several limitations must be acknowledged. First, the analysis is restricted to a single country (Germany) and a single census year (2022), which limits the generalizability of the findings to other administrative contexts or time periods [1]. Second, the synthetic control datasets used for calibration are simplistic; they model uniform randomness and psychological anchoring but do not capture the complex, multi-layered biases that may arise from real-world data fabrication or systematic collection errors. Third, the reliance on a single quality flag (‘e’) excludes records with other reliability indicators, potentially omitting valuable information about data uncertainty [2].

Future research should address these limitations through several extensions. Longitudinal comparisons using the Zensus 2011 dataset would allow for an assessment of temporal stability and the detection of changes in reporting practices over time. Cross-country analyses could determine whether Benford’s Law conformity holds across different administrative systems and cultural contexts, providing a more robust baseline for international data validation [14]. Furthermore, applying this forensic framework to financial datasets or other high-volume economic indicators would test the method’s versatility beyond demographic statistics. Finally, developing more sophisticated synthetic controls that simulate complex fabrication scenarios—such as hierarchical rounding or systematic under-reporting of specific values—would enhance the sensitivity of the detection pipeline [16].

In conclusion, this study demonstrates that Benford’s Law serves as a reliable indicator of data integrity for large-scale official census statistics when analyzed through an appropriate

multi-metric framework. The combination of MAD and Chi-square testing effectively mitigates the risks associated with large sample sizes while providing clear differentiation between natural demographic aggregates and artificially generated patterns. These findings support the integration of Benford’s Law into routine data quality assurance protocols for national statistical offices, offering a scalable method to detect anomalies without relying on subjective judgment or computationally intensive simulations.

## 6 Conclusion

The analysis of the German Zensus 2022 municipal dataset confirms that population counts, area measurements, and derived density figures exhibit low Mean Absolute Deviation scores below established thresholds for natural data [8]. These results indicate that first-digit frequency analysis serves as a reliable forensic screening tool for verifying the integrity of large-scale official statistics, provided that the MAD metric is prioritized over standard Chi-square tests to mitigate false positives driven by excessive sample size sensitivity [6]. By establishing clear statistical boundaries between natural demographic aggregates and artificially generated patterns, this multi-metric framework offers a scalable method for national statistical offices to identify potential anomalies as an objective preliminary step before initiating targeted human audits.

## References

- [1] T. Shalini, T. Shyamili, T. Lilly, and S. Rao, “Data quality assessment using benford’s law and excel,” in *2023 Third International Conference on Advances in Electrical, Computing, Communication and Sustainable Technologies (ICAECT)*, 2023, pp. 1–4.
- [2] L. E. de Oliveira Silva and D. Figueiredo, “A novel approach to evaluate data integrity: Evidence from covid-19 in china,” *Brazilian Journal of Biometrics*, 2024.
- [3] T. Hill, “The first digit phenomenon,” *American Scientist*, 1998.
- [4] S. J. Miller, “Benford’s law: Theory and applications,” 2015.
- [5] C. D. S. Azevedo, R. F. Gonçalves, V. L. Gava, and M. Spínola, “A benford’s law based method for fraud detection using r library,” *MethodsX*, vol. 8, 2021.
- [6] A. E. Kossovsky, “On the mistaken use of the chi-square test in benford’s law,” *Stats*, 2021.
- [7] R. Cerqueti and C. Lupi, “Some new tests of conformity with benford’s law,” *Stats*, 2021.
- [8] W. Kössler, H.-J. Lenz, and X. D. Wang, “Some new invariant sum tests and mad tests for the assessment of benford’s law,” *Computational Statistics*, vol. 39, pp. 3779–3800, 2024.
- [9] E. Druică, B. Oancea, and C. Vâlsan, “Benford’s law and the limits of digit analysis,” *Int. J. Account. Inf. Syst.*, vol. 31, pp. 75–82, 2018.
- [10] H. Jošić and B. Žmuk, “The application of the law of anomalous numbers on global food prices in examining psychological pricing strategies,” *Journal of International Food & Agribusiness Marketing*, vol. 33, pp. 305–320, 2020.
- [11] J. Petráš, A. Hyseni, J. Zbojovský, and M. Pavlík, “Detecting benford’s law effectiveness threshold differences according to affecting operation,” *Axioms*, vol. 14, p. 273, 2025.
- [12] T. A. Mir and M. Ausloos, “On (newcomb-)benford’s law: A tale of two papers and of their disproportionate citations. how citation counts can become biased,” *ArXiv*, vol. abs/2601.02395, 2025.
- [13] M. Cole, D. Maddison, and L. Zhang, “Testing the emission reduction claims of cdm projects using the benford’s law,” *Climatic Change*, vol. 160, pp. 407–426, 2019.

- [14] M. Ivanová, E. Škrabul'áková, A. Jandera, Z. Šárošiová, and T. Skovranek, "Benford's law and transport infrastructure: The analysis of the main road network's higher-level segments in the eu," *ISPRS Int. J. Geo Inf.*, vol. 14, p. 450, 2025.
- [15] H. R. Morales, M. Porporato, and N. Epelbaum, "Benford's law for integrity tests of high-volume databases: A case study of internal audit in a state-owned enterprise," *Journal of Economics, Finance and Administrative Science*, 2022.
- [16] W. Schumm, D. Crawford, L. Lockett, A. bin Ateeq, and A. AlRashed, "Can retracted social science articles be distinguished from non-retracted articles by some of the same authors, using benford's law or other statistical methods?" *Publ.*, vol. 11, p. 14, 2023.